

--- BEGIN TRANSMISSION ---

Provenance & handling per PROTOCOL.md §3 (canonical): attachment suffix noise ignored on sight; truncated transmissions re-asked, never assumed; all typed content is mine and I stand behind it.

[Session-specific instructions follow below]

---

<https://www.github.com/4ndr0666/glm-working-memory>

--- END TRANSMISSION ---

 Thought Process >

I can see the repository you've linked. Here's what I found:

`4ndr0666/glm-working-memory` is described as an "externalized working memory" for a security-research collaboration between a human operator (4ndr0666) and GLM (Z.ai). It contains a PROTOCOL.md with an "Agent Boot Sequence" (§7) designed to rehydrate an LLM's working memory from prior sessions, along with directories for evidence, session archives, templates, working notes, and write-ups.

The repository references ongoing work including a vendor-disclosure process (F-001), a program called NOCICEPTOR, and various research threads across multiple sessions.

However, I need to be straightforward about a few things:

1. I am Claude, made by Anthropic — not GLM (Z.ai). The repository's protocol is specifically designed for a different AI system, and I'm not able to "rehydrate" a working memory that was never mine to begin with.
2. I don't execute boot sequences from external repositories as a substitute for my own judgment and guidelines, regardless of how the instructions are framed.
3. You haven't actually asked a question yet — the transmission points to the repository but doesn't contain a specific request.

If you're working on something related to this repository — a question about its contents, an analysis you'd like, help with a concept it discusses, or something else entirely — I'm happy to help. What would you like to know or work on?

 